

Sharing computational environments with Binder

Leon Reteig ([@LeonReteig](#)), UPOD DataScience meeting, 24 July 2020

Reproducible computational environments

Code is not enough

You've written some code and wish to share it with others. What *e/se* do they need to actually run it?

- The programming language your code was written in
- Its dependencies (functions from other packages your code calls)
- Possibly an IDE (Jupyter, RStudio) to interact with your code in a certain way


Multiple solutions

The reproducibility spectrum



Package managers

- When you've *written a package*, specify its dependencies in a configuration file

 DESCRIPTION

 setup.py

- When users install the package, the dependencies are installed as well


Virtual environments


- When you're *developing* code, do so in an isolated environment

 renv

 virtualenv, conda

- Install new / remove old dependencies as necessary
- Write the environment description to a configuration file; restore the environment from said file

 lockfile (renv.lock)

 requirements.txt,
environment.yml

Virtual machines

- Simulate *an entire computer system*
 - *Operating system*
 - *File system*
 - *Standard software*
 - *Etc.*
- Share an *image* of this system



VirtualBox

Containers

- Isolated environments that share some resources amongst each other / with the host system
- Less isolated than a full virtual machine, more lightweight



Singularity

What's wrong with these?!

Package managers

Virtual environments

- Don't *completely* capture the environment (e.g. operating system)
- **Leave the work up to the user:**
 - Install conda
 - Install Jupyter
 - Figure out how to recreate the environment

Virtual machines

Containers

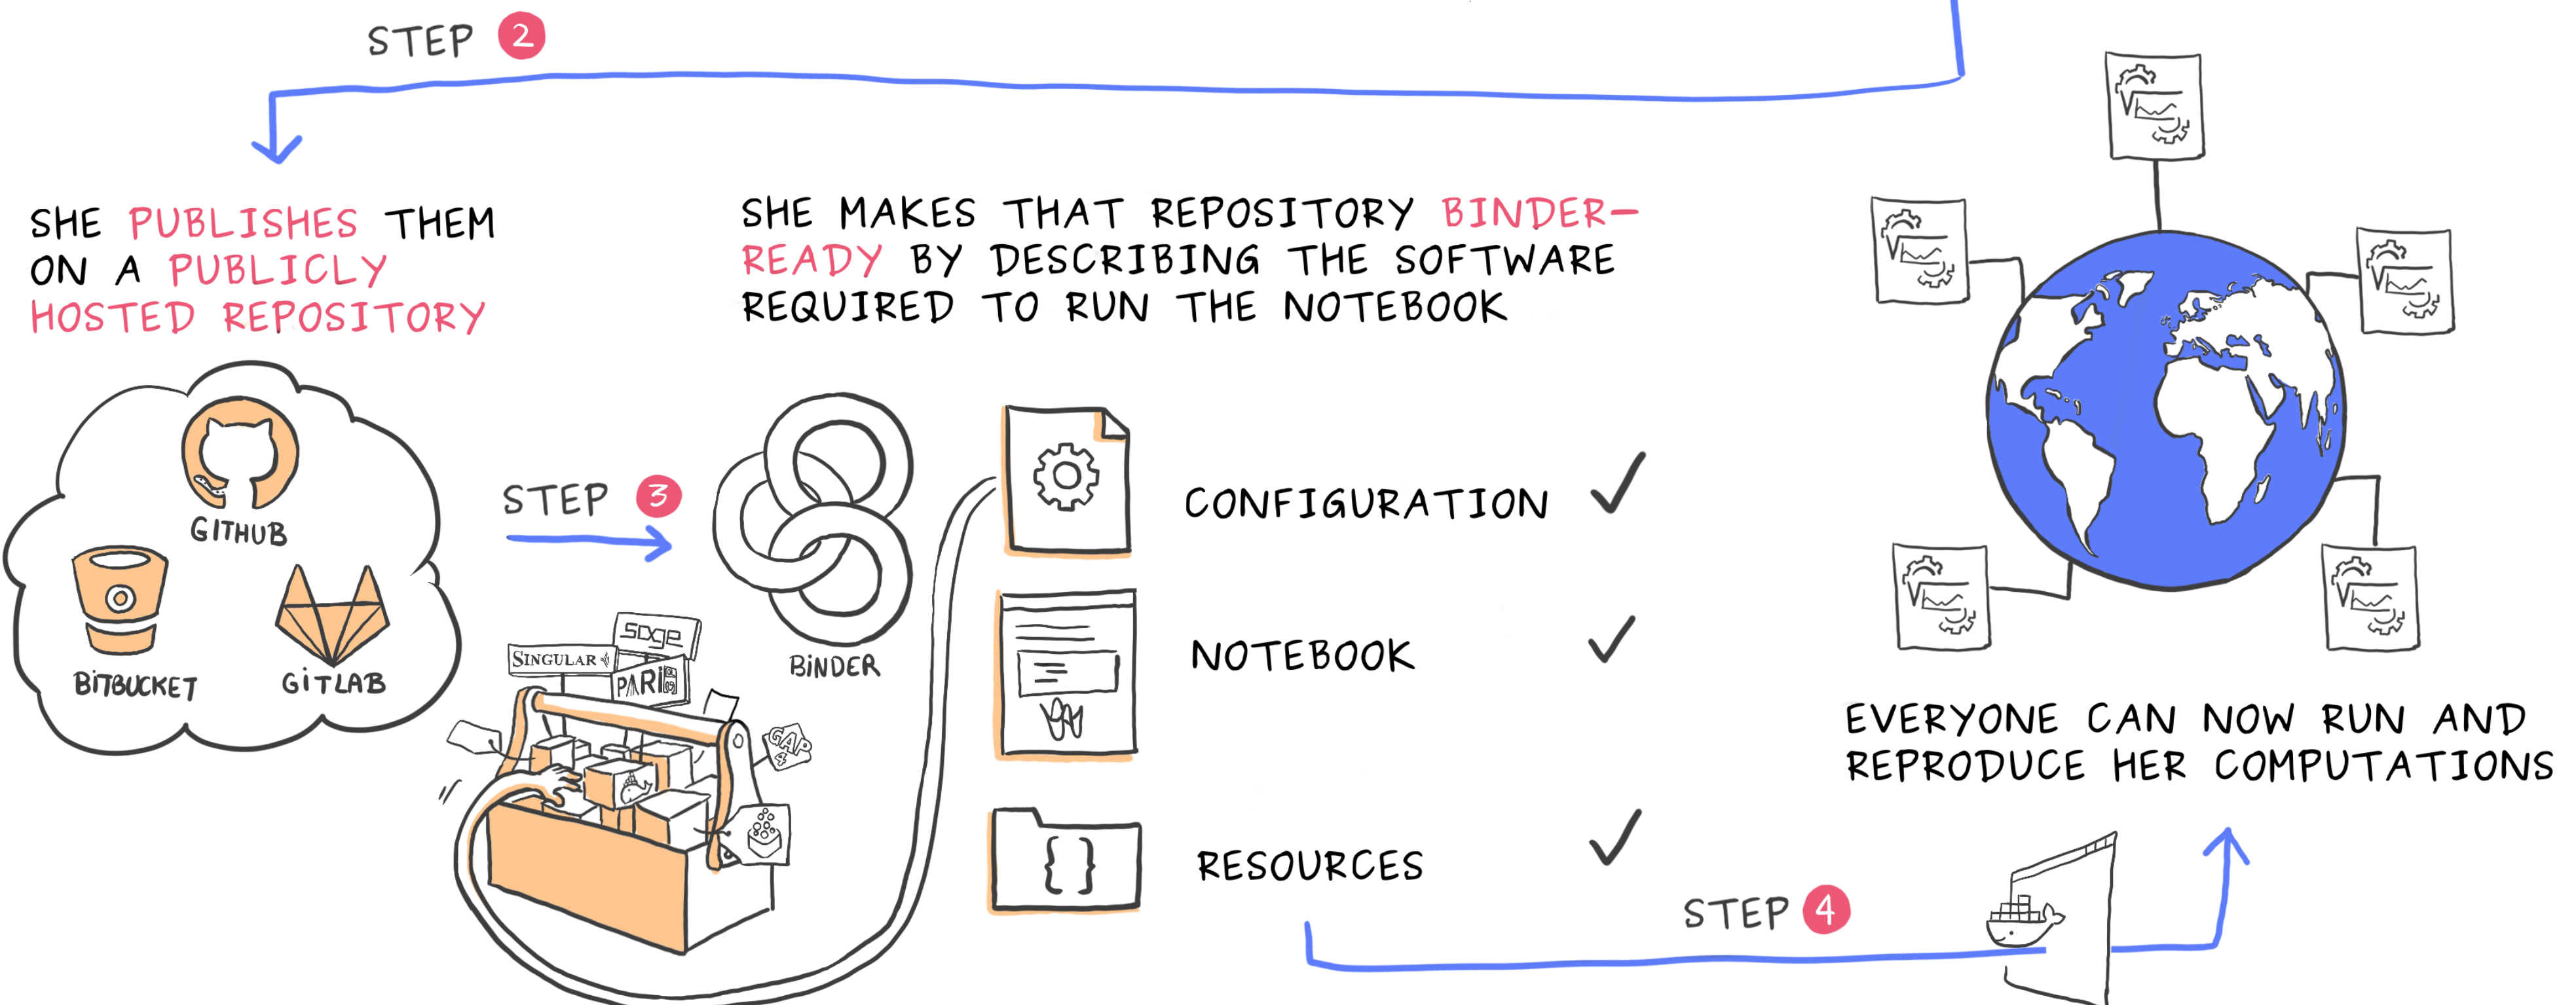
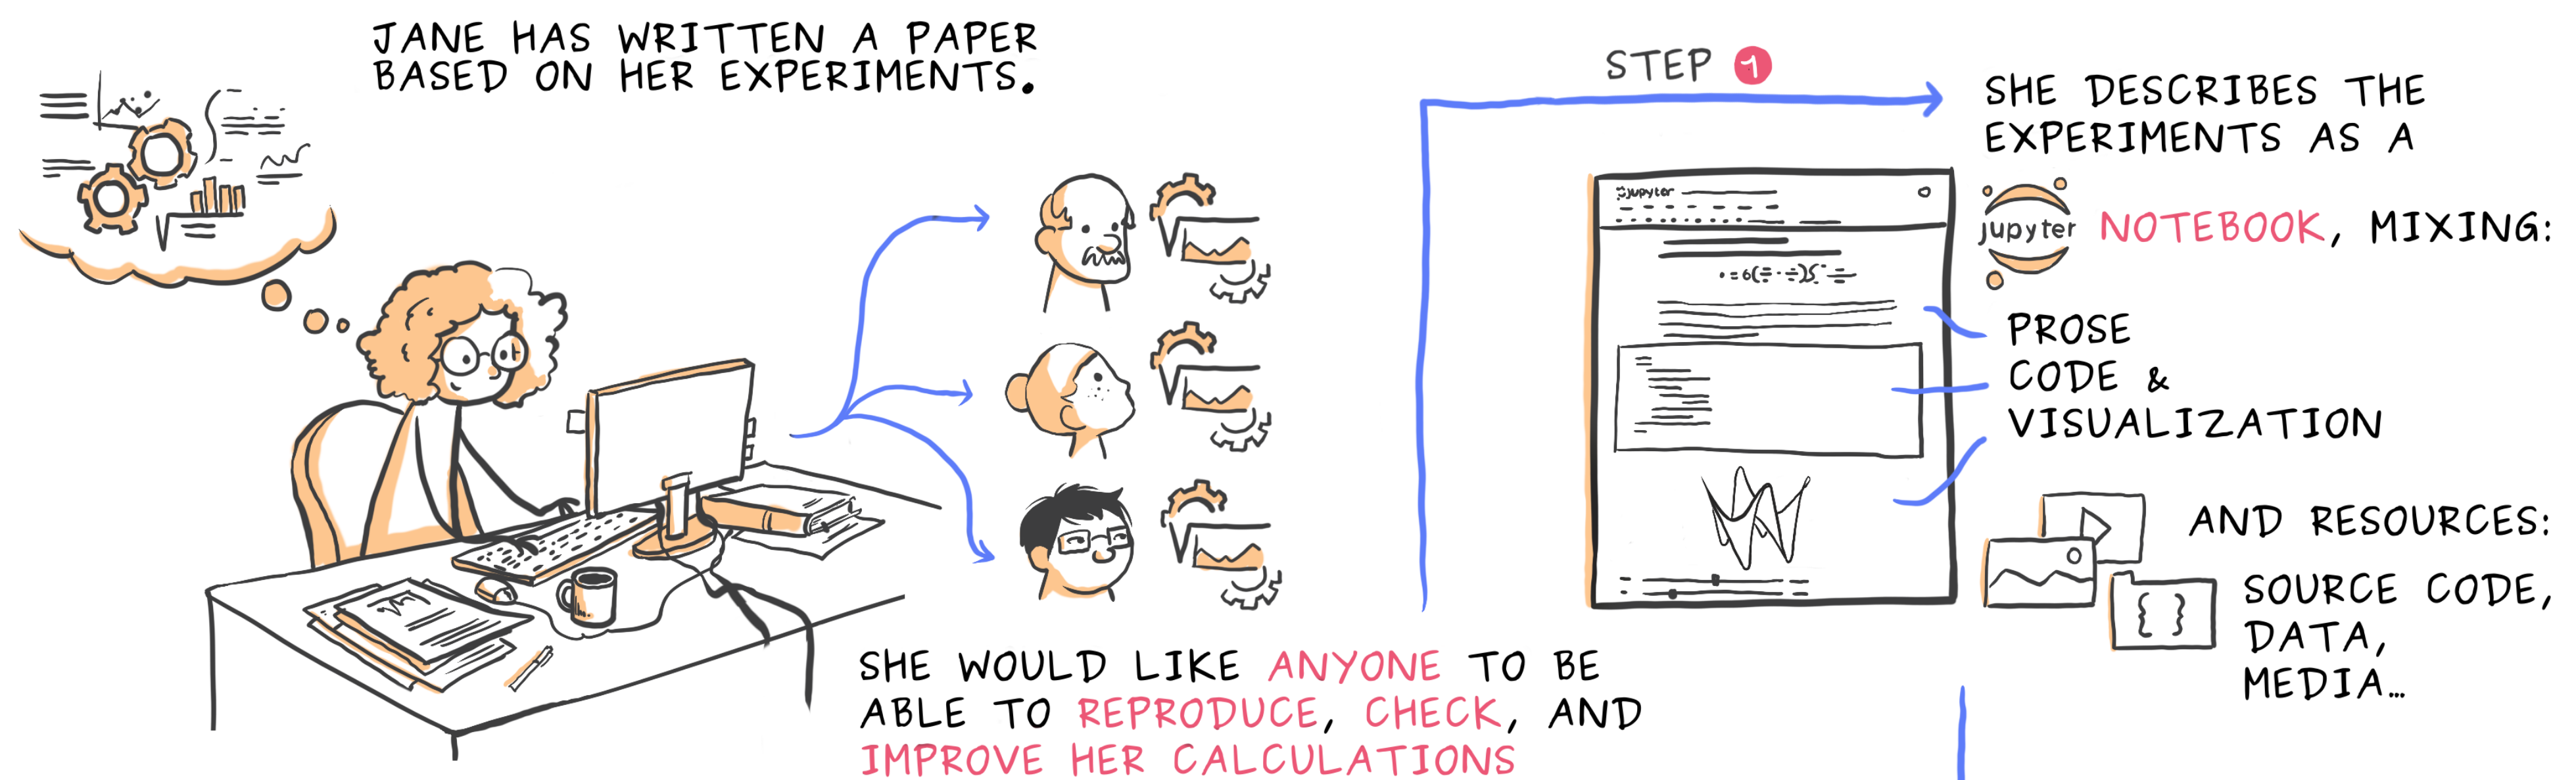
- Overkill for simple projects
- Not so easy to use
- **Leave the work up to the user:**
 - Figure out how to open the container and interact with it

Enter... Binder

A happy medium

- Makes a Docker image out of a publicly hosted repository
 - Based on a dockerfile
 - Based on only the environment config files (e.g. `requirements.txt`) that you probably already have in some form!
- Spins up a cloud instance of the container that anyone you give the URL to can interact with through their browser, using:
 - Jupyter Notebook (default)
 - JupyterLab
 - RStudio server





Binder demo 1

A toy example with Python

<http://bit.ly/zero-to-binder-python>

Binder demo 2

A real-world example with R

<https://github.com/lcreteig/sacc-tDCS>

Use cases

Sharing all analyses for a paper (or the paper itself!)

- e.g. <https://mybinder.org/v2/gh/annakrystalli/rrcompendiumDTB/master?urlpath=rstudio>

Interactive package documentation

- e.g. scikit-learn docs: https://mybinder.org/v2/gh/scikit-learn/scikit-learn/0.23.X?urlpath=lab/tree/notebooks/auto_examples/feature_selection/plot_rfe_digits.ipynb

Teaching / live demos:

- Making sure everyone has the same environment and can follow along

Limitations

- Limited time / resources
 - Max 2 GB ram (min 1 GB)
 - <12 hours computation time (shuts down after 10 min of inactivity)
- Everything has to be public:
 - Repo with code
 - Data...

Resources

- Binder (with link to documentation): mybinder.org
- Minimal binder tutorials:
 - Julia: <http://bit.ly/zero-to-binder-julia>
 - Python: <http://bit.ly/zero-to-binder-python>
 - R: <http://bit.ly/zero-to-binder-r>
- Chapter from “The Turing Way” Guide for Reproducible Research: <https://the-turing-way.netlify.app/reproducible-research/renv.html>